



# **Using regression and machine learning algorithm for housing price prediction in Ames, Iowa**

**Term paper of Econometrics II**

Ao Sun            18-624-007

Xiaoqian Zhu      18-624-056

May 5, 2019

## **Contents**

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Data</b>	<b>2</b>
<b>3</b>	<b>Methodology</b>	<b>3</b>
3.1	Identification strategy and estimation methods . . . . .	3
3.2	Advantages and disadvantages of the different estimators . . . . .	4
<b>4</b>	<b>Results</b>	<b>4</b>
4.1	Table of the results . . . . .	4
4.2	Compare the results of different estimators . . . . .	5
4.3	Critically discussion of the results. . . . .	5
<b>5</b>	<b>Conclusions</b>	<b>7</b>
	<b>Reference</b>	<b>8</b>

# 1 Introduction

Real estate industry is the basic guarantee of people's life, and houses is one of the most common way to accumulating wealth. In the last decade, housing prices in most cities have been on the rise. Whether the prices will continue to rise in the future, and to what level they can reach, it's highly concerned among the government, developers, and ordinary citizens.

Liu (2014) pointed that accurate forecasting for future housing price is very important for socioeconomic development and national lives and helps in establishing real estate policies.

In this paper we used four different methods to fit the data, OLS, LASSO and Ridge regression and Random Forest. After that, we used bootstrap to get the performance of the estimators under each model. It turns out that Random forest is superior to the others on both test and train data.

The analysis of housing price has two principal trends: one is called hedonic-based regression, which is more concentrating on the market fundamentals.

Bin(2004) used a semi-parametric method to estimate the hedonic price function and find it performances better than conventional parametric model. There are lots of studies use machine learning method to do the prediction. Park and Bae(2015) developed a prediction model of the Fairfax County of Virgina based on naïve Bayesian, Adaboost methods, and found that ML algorithm enhance the predictability and significantly contribute on the correct evaluation of real estate price. Gu, Zhu and Jiang(2011) suggested a hybrid model of genetic algorithm and support vector machine(G-SVM) model to predict the housing price and found that this model has higher forecast accuracy.

# 2 Data

We get 79 explanatory variables describing almost every aspect of residential homes with the *SalePrice* in Ames of Iowa from Kaggle Competition. After the visualization of the *SalePrice*, we found that it fits the log-normal distribution well.




Figure 1: Distribution visualization of Sale Price

There are several missing values, we drop those variables with missing values more than 10% of the number observations. And transformed non-numerical

labels (as long as they are comparable) to numerical labels.

From the plot we can observe outliers, we delete several unreasonable outliers of variables *GrLivArea* and *TotalBsmtSF*.

In order to compare the performance of each model, we split 20% of the observations as test data, and do the fitting on the rest (train) data.

Until now still remain many explanatory variables, so we get the correlation matrix and use *SelectKBest* in python to get the first fifteen important features. The correlations as follows:




Figure 2: Correlation matrix

### 3 Methodology

#### 3.1 Identification strategy and estimation methods

In this paper, we selected four estimation methods. They are OLS, Ridge(Hoerl, Kennard, 1970), Lasso(Tibshirani, 1996) and Random Forest (Breiman, 2001) respectively. Since the first three model have linear expressions, we do the variance comparison via bootstrap between those three models.

Before fitting the model, we had selected 15 explanatory variables including *OverallQual*, *YearBuilt*, *YearRemodAdd* etc, to predict the log *SalePrice* of house. OLS, Ridge and Lasso are all in generalized linear model classes, which means we can get the estimated weight parameters  $\hat{\beta}$  after fitting the models. While Random Forest is a non-linear model, from which we get the variables importance instead of the estimated parameters.

We use Grid Search (Bergstra, Bengio, 2012) method to select hyper parameters during fitting process.

After the model fitting, we discussed the performance of the models in test data. We chose MSE as the criterion and do bootstrap to estimate the variance of  $\hat{\beta}$  in linear models and variables importance (Mean decrease impurity) in Random Forest.

### 3.2 Advantages and disadvantages of the different estimators

In our data, *SalePrice* is a continuous variable, there are two advantages of OLS.

Firstly, it's a simple method to explain and to understand.  $\hat{\beta}$  suggests how many *SalePrice* will change with one unit change of explanatory variables.

Secondly, it's applicable in most general condition whether the *SalePrice* is continuous or discrete. If *SalePrice* is a binary variable, perhaps use probit or logit model will be better than classical OLS. But OLS model can be used in most regression tasks.

Compared to other models, the disadvantages of OLS are main in three points. Primarily, it's sensitive to outliers, so that's why we've removed the outliers in training data. Then, when the data is not normally distributed, the outcome might be unreliable, so we've used a relatively big data set to alleviate the problem. In addition, OLS has tendency to overfit data.

Lasso and Ridge can help us solve the last disadvantage of OLS. In Lasso we can use L1 regularization to shrinking some elements in  $\beta$  to 0. That means Lasso is useful for variables selection by which the redundant variables are discarded, while Ridge use L2 regularization to make the redundant elements relatively small. Including more variables means larger probability to meet the overfitting problem, since they can help decrease the effects of redundant variables, they become two common way to avoid overfitting. The most disadvantages of Lasso and Ridge is the bigger computation time than OLS because of the regularization.

The last selected model is Random Forest. The biggest difference is that Random Forest is a non-linear model. It has different criterion compare to linear model. Breiman (2001) pointed out some advantages of Random Forest:

Firstly, using random features can accelerate the model fitting process.

Next, Random Forest is less sensitive to outliers than OLS.

Moreover, Random Forest is an ensemble model which combine many decision trees and use bagging methods to avoid overfitting.

The biggest disadvantage of Random Forest is the curse of dimensionality. And it doesn't give much economics explanation as the above three models do.

## 4 Results

### 4.1 Table of the results

After the point estimation, we did 100 times bootstrap for variance estimation. Since Efron(1993) pointed that 50-200 times resampling is enough for variance estimation. We got result tables for all models:

	Random Forest		OLS		Ridge		Lasso	
	Mean decrease impurity	std	$\hat{\beta}$	std	$\hat{\beta}$	std	$\hat{\beta}$	std
OverallQual	182.42	11.74	113.90	9.39	107.08	8.96	108.84	8.91
YearBuilt	102.22	13.09	1.93	0.39	1.94	0.38	1.92	0.37
YearRemodAdd	40.86	9.27	3.36	0.43	3.43	0.43	3.45	0.43
ExterQual	63.59	12.62	16.37	14.87	10.74	12.77	<b>0.00</b>	4.07
BsmtQual	29.46	8.29	-15.41	8.56	-15.30	8.06	-10.16	7.52
TotalBsmtSF	71.64	9.91	0.24	0.03	0.24	0.03	0.24	0.03
1stFlrSF	62.50	7.83	0.06	0.04	0.06	0.04	0.06	0.04
GrLivArea	155.99	13.87	0.39	0.03	0.39	0.03	0.37	0.02
FullBath	52.93	11.28	-34.32	19.27	-25.81	15.21	<b>0.00</b>	12.44
KitchenQual	26.46	8.49	-25.71	9.62	-23.68	8.63	-12.87	9.61
TotRmsAbvGrd	22.09	3.03	-6.92	7.85	-7.31	7.51	-3.07	5.65
GarageType	32.14	11.07	-11.50	4.49	-11.87	4.44	-10.46	4.42
GarageCars	69.48	9.73	26.57	22.75	20.01	16.57	<b>0.00</b>	9.10
GarageArea	61.55	8.81	0.23	0.07	0.26	0.05	0.32	0.04
MeanArea	26.67	4.12	0.12	0.03	0.12	0.03	0.12	0.02
MSE in out of sample data	51.54	1.97	104.55	4.57	105.90	4.37	107.45	4.12

\*All number have been multiplied with 1000 since the values are extremely small.

Table 1: Table of results

## 4.2 Compare the results of different estimators

Firstly, we compared the results between linear models. And then compare the best model among linear models with Random Forest.

1. OLS has the minimum MSE in test data. Second is Ridge. Lasso performances the worst among linear models based on this criterion.
2. *FullBath* has the biggest negative value in OLS and Ridge, but in Lasso, the L1 regularization make it zero.
3. Most variables' values in Ridge are less than those in OLS.
4. The standard deviations in OLS are always bigger than those in Ridge and Lasso. But between Ridge and Lasso, the standard deviations have no obvious difference.
5. *ExterQual*, *FullBath*, *GarageCars* have been assigned to 0 in Lasso, which means Lasso model think they are redundant variables.

Then compare the Random Forest and linear models:

1. Random Forest gets better performance on test data than linear models.
2. Random Forest considers *OverallQual* has the biggest influence, too. But Random Forest think the next important variable is *GrLivArea* which has a relatively small weight in linear model.

## 4.3 Critically discussion of the results.

As what we've concluded above, OLS model has the biggest standard deviation among linear models in train data. Even OLS got the best out sample performance in linear models, we can't conclude the OLS is the best model for prediction still. In addition, Lasso model discarded *ExterQual* (Evaluates the quality of the material on the exterior), *FullBath* (Full bathrooms above grade) and *GarageCars* (Size of garage in

car capacity). Although Lasso got the smallest standard deviations, the discarded variables can intuitively affect house price much in reality. We thought this might be the main reason that Lasso got the worst performance in out sample.

Besides, the MSE of Random Forest in test data is only half of the linear models. Even we used non-traditional linear model, like Ridge and Lasso, Random forests still outperform. That may mean linear models have more limitations to explain the explanatory ability of variables on *SalePrice*.

We proposed 3 reason why linear model did much worse than Random Forest:

1. OLS, Lasso and Ridge are one-dimensional linear models, they are much more simpler.
2. Many variables are discrete factors rather than continuous number. For example, variable *OverallQual*(Rates the overall material and finish of the house) was assigned as integers from 1 to 10 which is more effective in Random Forest estimation.
3. The outliers affected linear models more.




Figure 3: Comparison of test data with predicted data

From the above figure, we can see that there are some outliers in the prediction of linear model. However, outliers affected Random Forest prediction less than other models.

We thought the results is plausible. In all models, variable *OverallQual* is the primary factor affecting the house price. That outcome is in accord with our experience in reality. Moreover, in Random Forest model, the area variables affect housing price a lot and so do the linear model. Considering the lager value scale measurement of area

variables, we think the result is intuitive.

## 5 Conclusions

From the analysis above, we found Random Forest perform much better than linear models under out of sample MSE criterion.

Among linear models, OLS had the lowest out of sample MSE, but it had the biggest variance of estimated parameters  $\hat{\beta}$ . Lasso got the lowest variance, but because of discarding some variables, it had the largest out of sample MSE. Compared to OLS and Lasso, Ridge got an intermediate performance on either out of sample MSE or the variance of estimated parameters.

We think the Random Forest is the best model to do the prediction task, and we proposed three ways that probably can help increase the out of sample performance:

1. Use another criterion to select variables and test different variables combination, like information entropy.
2. Do normalization and standardization on some variables to align the scale of variables.
3. Check and remove more outliers in the training data, since we've got a lot explanatory variables.
4. Most literature now use suggest to do the Box-Cox transformation, which can be treated as a generalized log transformation, and most researchers also include higher-order terms in their regressions. Those methods could also be used for improving our model.

## References

- Bin, O. (2004). A prediction comparison of housing sales prices by parametric versus semi-parametric regressions. *Journal of Housing Economics*. 13(1), 68-84.
- Bergstra, J., & Bengio, Y. (2012). Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(1), 281-305.
- Bischoff, E., Bischoff, H., & Giuliano, F. (2005). An introduction to the bootstrap / b. efron, r.j. tibshirani.
- Friedman, J. H. . (2001). Greedy function approximation: a gradient boosting machine. *The Annals of Statistics* , 29(5), 1189-1232.
- Gu, J., Zhu, M., & Jiang, L. (2011). Housing price forecasting based on genetic algorithm and support vector machine. *Expert Systems with Applications* , 38(4), 3383-3386.
- Hoerl, A., & Kennard, R. (1970). Ridge regression: applications to nonorthogonal problems. *Technometrics*, 12(1), 14.
- Liu, J. G. , Zhang, X. L. , & Wu, W. P. . (2006). Application of fuzzy neural network for real estate prediction. *International Symposium on Advances in Neural Networks-isnn*. DBLP
- Park, B., & Bae, J. K. (2015). Using machine learning algorithms for housing price prediction: the case of fairfax county, virginia housing data. *Expert Systems with Applications* , 42(6), 2928-2934.
- Tibshirani, R. (1996). Regression shrinkage and selection via the lasso: a retrospective. *Journal of the Royal Statistical Society* , 58(1), 267-288.

# Appendix

May 5, 2019

## 0.1 Import libraries

```
In [1]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import seaborn as sns
        from scipy import stats
        from scipy.stats import norm
        from sklearn.preprocessing import LabelEncoder
        from sklearn.ensemble import RandomForestRegressor
        from sklearn.linear_model import LinearRegression, Lasso, Ridge
        from sklearn.metrics import mean_squared_error
        from sklearn.svm import SVR
        from sklearn.feature_selection import SelectKBest
        from sklearn.feature_selection import f_regression
        from sklearn.model_selection import GridSearchCV
        from sklearn.utils import resample
```

```
In [2]: ##### set the random seed
        np.random.seed(137)
```

```
In [3]: color = sns.color_palette()
        sns.set_style('darkgrid')
```

## 1 1. Import data

```
In [4]: # original_data = pd.read_csv('houspr.csv', index_col=0).drop(['Alley', 'PoolQC', 'Fen
        original_data = pd.read_csv('houspr.csv', index_col=0)
        train_data = original_data.iloc[:1200]
        test_data = original_data.iloc[1200:]
```

```
In [5]: original_data.head()
```

```
Out[5]:    MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
Id
1          60      RL       65.0     8450    Pave    NaN    Reg
2          20      RL       80.0     9600    Pave    NaN    Reg
3          60      RL       68.0    11250    Pave    NaN   IR1
```

4	70	RL	60.0	9550	Pave	NaN	IR1	
5	60	RL	84.0	14260	Pave	NaN	IR1	
LandContour Utilities LotConfig				...	PoolArea	PoolQC	Fence	\
Id				...				
1	Lvl	AllPub	Inside	...	0	NaN	NaN	
2	Lvl	AllPub	FR2	...	0	NaN	NaN	
3	Lvl	AllPub	Inside	...	0	NaN	NaN	
4	Lvl	AllPub	Corner	...	0	NaN	NaN	
5	Lvl	AllPub	FR2	...	0	NaN	NaN	
MiscFeature MiscVal MoSold YrSold SaleType SaleCondition SalePrice								
Id								
1	NaN	0	2	2008	WD	Normal	208500	
2	NaN	0	5	2007	WD	Normal	181500	
3	NaN	0	9	2008	WD	Normal	223500	
4	NaN	0	2	2006	WD	Abnorml	140000	
5	NaN	0	12	2008	WD	Normal	250000	

[5 rows x 80 columns]

### 1.0.1 statistics of the data

In [6]: original\_data.dtypes

Out [6]:

MSSubClass	int64
MSZoning	object
LotFrontage	float64
LotArea	int64
Street	object
Alley	object
LotShape	object
LandContour	object
Utilities	object
LotConfig	object
LandSlope	object
Neighborhood	object
Condition1	object
Condition2	object
BldgType	object
HouseStyle	object
OverallQual	int64
OverallCond	int64
YearBuilt	int64
YearRemodAdd	int64
RoofStyle	object
RoofMatl	object
Exterior1st	object

```

Exterior2nd      object
MasVnrType       object
MasVnrArea        float64
ExterQual         object
ExterCond         object
Foundation        object
BsmtQual          object
...
BedroomAbvGr     int64
KitchenAbvGr     int64
KitchenQual       object
TotRmsAbvGrd     int64
Functional        object
Fireplaces        int64
FireplaceQu      object
GarageType        object
GarageYrBlt      float64
GarageFinish      object
GarageCars        int64
GarageArea        int64
GarageQual        object
GarageCond        object
PavedDrive        object
WoodDeckSF        int64
OpenPorchSF       int64
EnclosedPorch     int64
3SsnPorch         int64
ScreenPorch       int64
PoolArea          int64
PoolQC            object
Fence              object
MiscFeature       object
MiscVal            int64
MoSold             int64
YrSold             int64
SaleType           object
SaleCondition      object
SalePrice          int64
Length: 80, dtype: object

```

In [7]: `original_data.describe(include='all')`

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	\
count	1460.000000	1460	1201.000000	1460.000000	1460	91	
unique	NaN	5	NaN	NaN	2	2	
top	NaN	RL	NaN	NaN	Pave	Grvl	
freq	NaN	1151	NaN	NaN	1454	50	
mean	56.897260	NaN	70.049958	10516.828082	NaN	NaN	

	LotShape	LandContour	Utilities	LotConfig	...	PoolArea	\
count	1460	1460	1460	1460	...	1460.000000	
unique	4	4	2	5	...		NaN
top	Reg	Lvl	AllPub	Inside	...		NaN
freq	925	1311	1459	1052	...		NaN
mean	NaN	NaN	NaN	NaN	...	2.758904	
std	NaN	NaN	NaN	NaN	...	40.177307	
min	NaN	NaN	NaN	NaN	...	0.000000	
25%	NaN	NaN	NaN	NaN	...	0.000000	
50%	NaN	NaN	NaN	NaN	...	0.000000	
75%	NaN	NaN	NaN	NaN	...	0.000000	
max	NaN	NaN	NaN	NaN	...	738.000000	
	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	\
count	7	281	54	1460.000000	1460.000000	1460.000000	
unique	3	4	4	NaN	NaN	NaN	
top	Gd	MnPrv	Shed	NaN	NaN	NaN	
freq	3	157	49	NaN	NaN	NaN	
mean	NaN	NaN	NaN	43.489041	6.321918	2007.815753	
std	NaN	NaN	NaN	496.123024	2.703626	1.328095	
min	NaN	NaN	NaN	0.000000	1.000000	2006.000000	
25%	NaN	NaN	NaN	0.000000	5.000000	2007.000000	
50%	NaN	NaN	NaN	0.000000	6.000000	2008.000000	
75%	NaN	NaN	NaN	0.000000	8.000000	2009.000000	
max	NaN	NaN	NaN	15500.000000	12.000000	2010.000000	
	SaleType	SaleCondition	SalePrice				
count	1460	1460	1460.000000				
unique	9	6	NaN				
top	WD	Normal	NaN				
freq	1267	1198	NaN				
mean	NaN	NaN	180921.195890				
std	NaN	NaN	79442.502883				
min	NaN	NaN	34900.000000				
25%	NaN	NaN	129975.000000				
50%	NaN	NaN	163000.000000				
75%	NaN	NaN	214000.000000				
max	NaN	NaN	755000.000000				

[11 rows x 80 columns]

### 1.0.2 1.1 Fit normal distribution with original data(sale\_price) and compare

```
In [8]: # sns.distplot(train_data['SalePrice'], fit=norm)
(mu, sigma) = norm.fit(train_data['SalePrice'])
print("The mean of normal fitting the orginal data is {} \nThe standard deviation of no
```

The mean of normal fitting the orginal data is 181414.6283

The standard deviation of normal fitting the orginal data is 81037.122


### 1.0.3 1.2 Plot the distribution of sale\_price(original data)


```
In [54]: plt.figure(figsize=(14, 5))
sns.distplot(train_data['SalePrice'], fit=norm)
plt.savefig("normal distribution.jpg")
```

```
# Get the QQ-plot
plt.figure(figsize=(14, 5))
res = stats.probplot(train_data['SalePrice'], plot=plt)

plt.show()
```

```
/anaconda3/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-t
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```





## 2 2. Data processing

### 2.1 2.1 Do the log transformation on sale\_price

```
In [56]: train_data.loc[:, "SalePrice"] = np.log1p(train_data["SalePrice"])
```

#### 2.1.1 2.1.1 Fit normal distribution with the data after log-transformation and compare

```
In [57]: (mu1, sigma1) = norm.fit(train_data['SalePrice'])
```

```
print("The mean of normal fitting the log-data is {} \nThe standard deviation of normal fitting the log-data is {}")
```

The mean of normal fitting the log-data is 12.0243

The standard deviation of normal fitting the log-data is 0.4034

#### 2.1.2 2.1.2 Plot the distribution of log sale\_price

```
In [58]: plt.figure(figsize=(14, 5))
```

```
sns.distplot(train_data['SalePrice'], fit=norm)
```


```
plt.savefig("log-normal distribution.jpg")
```

```
# Get also the QQ-plot
```

```
plt.figure(figsize=(14, 5))
```

```
res = stats.probplot(train_data['SalePrice'], plot=plt)
```


```
plt.show()
```



## 2.2 2.2 Plot the correlation matrix

```
In [19]: plt.subplots(figsize=(9, 9))
corr_x = train_data.corr()
sns.heatmap(corr_x, annot=False, vmax=1, square=True, cmap="Blues")
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1a19631be0>
```




### 2.3 2.3 Dealing with the missing values

```
In [20]: missing_pd = (original_data.isnull().sum() / original_data.shape[0] * 100).sort_values()
missing_pd.head(20)

f, ax = plt.subplots(figsize=(15, 12))
plt.xticks(rotation='90')
sns.barplot(x=missing_pd.index, y=missing_pd)
plt.xlabel('Features', fontsize=15)
plt.ylabel('Percent of missing values', fontsize=15)
plt.title('Percent missing data by feature', fontsize=15)

Out[20]: Text(0.5, 1.0, 'Percent missing data by feature')
```



## 2.4 2.4 we drop the variables which missing data ratio exceeding 10%

```
In [21]: original_data = original_data.drop(missing_pd[missing_pd > 10].index, axis=1)
```

```
missing_pd = (original_data.isnull().sum() / original_data.shape[0] * 100).sort_values()
```

```
missing_variables = missing_pd[missing_pd > 0].index
```

```
In [22]: # we still have 13 variables to deal with
```

```
# Replace GarageType, GarageFinish, GarageQual and GarageCond missing value with None
```

```
for var in ('GarageType', 'GarageFinish', 'GarageQual', 'GarageCond'):
    original_data[var] = original_data[var].fillna('None')
```

```
In [23]: # Replace GarageYrBlt missing value with 0
```

```
original_data['GarageYrBlt'] = original_data['GarageYrBlt'].fillna(0)
```

```
In [24]: # Replace BsmtFinType2, BsmtExposure, BsmtQual, BsmtFinType1, BsmtCond missing value with None
```

```
for var in ('BsmtFinType2', 'BsmtExposure', 'BsmtQual', 'BsmtFinType1', 'BsmtCond'):
    original_data[var] = original_data[var].fillna('None')
```

```
In [25]: # Replace MasVnrArea missing with 0
original_data["MasVnrArea"] = original_data["MasVnrArea"].fillna(0)

In [26]: # Replace MasVnrType missing with None
original_data["MasVnrType"] = original_data["MasVnrType"].fillna("None")

In [27]: # Replace Electrical missing with most frequent value
original_data['Electrical'] = original_data['Electrical'].fillna(original_data['Electrical'].mode()[0])
```

## 2.4.1 check the missing value

```
In [28]: missing_pd = (original_data.isnull().sum() / original_data.shape[0] * 100).sort_values(ascending=True)
missing_variables = missing_pd[missing_pd > 0].index
print(missing_variables.shape)

(0,)
```

## 2.5 2.5 Now we have 73 variables

### 2.5.1 2.5.1 Transform non-numerical labels to numerical labels

```
In [29]: original_data['MSSubClass'] = original_data['MSSubClass'].astype('str')
original_data['OverallCond'] = original_data['MSSubClass'].astype('str')
original_data['YrSold'] = original_data['YrSold'].astype(str)
original_data['MoSold'] = original_data['MoSold'].astype(str)

### Label Encoding
s_variables = original_data.columns[(original_data.dtypes == 'object')]
for col in s_variables:
    lbl = LabelEncoder()
    lbl.fit(list(original_data[col].values))
    original_data[col] = lbl.transform(list(original_data[col].values))
```

## 2.6 2.6 features engineering

```
In [30]: original_data['MeanArea'] = original_data['MSSubClass'].map(
    lambda x: original_data.groupby('MSSubClass')[['GrLivArea', 'TotalBsmtSF']].mean().sum(1))
original_data['MeanGrLivArea'] = original_data['MSSubClass'].map(
    lambda x: original_data.groupby('MSSubClass')[['GrLivArea']].mean())
original_data['MeanTotalBsmtSF'] = original_data['MSSubClass'].map(
    lambda x: original_data.groupby('MSSubClass')[['TotalBsmtSF']].mean())

In [31]: # generate the log sale_price
original_data['LogSalePrice'] = np.log2(original_data['SalePrice'])

In [32]: # inspect the transformed data
train_data = original_data.iloc[:1200]

# correlation matrix
```


```

plt.subplots(figsize=(12, 12))
corr_x = train_data.corr()
sns.heatmap(corr_x, annot=False, vmax=1, square=True, cmap="Blues")

corr_x = train_data.corr()
# select the variables that have correlation greater than 0.3
variables = train_data.columns[abs(corr_x.LogSalePrice) > 0.4]

# extract the 21 variables data
train_data = train_data[variables]
test_data = original_data.iloc[1200:][variables]

```



## 2.7 2.7 Explore the outliers

```
In [33]: fig, ax = plt.subplots(figsize=(9, 9))
```

```
    ax.scatter(x=train_data['GrLivArea'], y=train_data['SalePrice'])
    plt.ylabel('SalePrice', fontsize=13)
    plt.xlabel('GrLivArea', fontsize=13)
    plt.show()
```



### 2.7.1 2.7.1 Deleting outliers

```
In [36]: train_data = train_data.drop(train_data[(train_data['GrLivArea'] > 4000) & (train_data['SalePrice'] > 700000)])
# Check the graphic again
fig, ax = plt.subplots(figsize=(9, 9))
ax.scatter(train_data['GrLivArea'], train_data['SalePrice'])
```

```
plt.ylabel('SalePrice', fontsize=13)
plt.xlabel('GrLivArea', fontsize=13)
plt.show()
```



## 2.7.2 Explore another outliers

```
In [37]: fig, ax = plt.subplots(figsize=(9, 9))
ax.scatter(x=train_data['TotalBsmtSF'], y=train_data['SalePrice'])
plt.ylabel('SalePrice', fontsize=13)
plt.xlabel('TotalBsmtSF', fontsize=13)
plt.show()
```



### 2.7.3 Deleting outliers

```
In [38]: train_data = train_data.drop(train_data[(train_data['TotalBsmtSF'] > 3000) & (train_da
```

```
In [39]: # Check the graphic again
fig, ax = plt.subplots(figsize=(9, 9))
ax.scatter(train_data['TotalBsmtSF'], train_data['SalePrice'])
plt.ylabel('SalePrice', fontsize=13)
plt.xlabel('TotalBsmtSF', fontsize=13)
plt.show()
```



## 2.8 2.8 Get training data and test data

```
In [40]: train_y = train_data['LogSalePrice']
train_x = train_data.drop(['LogSalePrice', 'SalePrice'], axis=1)

test_y = test_data['LogSalePrice']
test_x = test_data.drop(['LogSalePrice', 'SalePrice'], axis=1)
```

## 3 3. Features selection

```
In [81]: X_new = SelectKBest(f_regression, k=15).fit_transform(train_x, train_y)
features = []
j = 0
for i in range(train_x.shape[1]):
```

```

if all(X_new[:, j] == train_x.iloc[:, i]):
    features.append(train_x.columns[i])
    j += 1
if j == 15:
    break
##### now we have 15 features


```

In [84]:

```

df = pd.DataFrame(X_new)
plt.subplots(figsize=(9, 9))
correlations = df.corr()
sns.heatmap(correlations, annot=False, vmax=1, square=True, cmap="Blues")
plt.savefig("correlation.jpg")
plt.show()

```



## 4 4. Fit the models

```
In [42]: train_x = train_x[features]
         test_x = test_x[features]

         outcome = {
             'RF': pd.DataFrame(),
             'OLS': pd.DataFrame(),
             'RIDGE': pd.DataFrame(),
             'LASSO': pd.DataFrame()
         }

In [43]: def models(train_x, train_y):

    #####
    ##### RandomForest #####
    #####
    # Random Forest Regressor parameters
    rf_params = {
        'n_jobs': -1,
        'n_estimators': 300,
        'warm_start': True,
        'max_depth': 15,
        'min_samples_leaf': 9,
        'max_features': 'log2',
        'verbose': 0
    }
    model = GridSearchCV(RandomForestRegressor(**rf_params),
                          {'max_depth': [10, 12, 14],
                           'min_samples_leaf': [1, 2]},
                          n_jobs=-1, cv=5)
    model_1 = model.fit(train_x, train_y)

    model = RandomForestRegressor(**rf_params)
    model = model.set_params(**model_1.best_params_)
    model_1 = model.fit(train_x, train_y)

    m1 = mean_squared_error(test_y, model_1.predict(test_x))

    df_1 = pd.Series(model_1.feature_importances_, index=train_x.columns)
    df_1['mean_squared_error'] = m1

    #####
    ##### OLS #####
    #####
    ols_params = {
        'copy_X': True,
        'fit_intercept': True,
```

```

        'n_jobs': None,
        'normalize': False
    }
model = LinearRegression(**ols_params)
model_2 = model.fit(train_x, train_y)

m2 = mean_squared_error(test_y, model_2.predict(test_x))

df_2 = pd.Series(model_2.coef_, index=train_x.columns)
df_2['mean_squared_error'] = m2

#####
##### Ridge #####
#####
rid_params = {
    'alpha': 0.3,
    'copy_X': True,
    'fit_intercept': True,
    'normalize': False
}
model = GridSearchCV(Ridge(**rid_params),
                     {'alpha': [40, 50, 60]},
                     n_jobs=-1, cv=5)
model_3 = model.fit(train_x, train_y)

model = Ridge(**rid_params)
model = model.set_params(**model_3.best_params_)
model_3 = model.fit(train_x, train_y)

m3 = mean_squared_error(test_y, model_3.predict(test_x))

df_3 = pd.Series(model_3.coef_, index=train_x.columns)
df_3['mean_squared_error'] = m3

#####
##### Lasso #####
#####
las_params = {
    'alpha': 0.1,
    'copy_X': True,
    'fit_intercept': True,
    'normalize': False
}
model = GridSearchCV(Lasso(**las_params),
                     {'alpha': [0.004, 0.005, 0.006]},
                     n_jobs=-1, cv=5)
model_4 = model.fit(train_x, train_y)

```

```

model = Lasso(**las_params)
model = model.set_params(**model_4.best_params_)
model_4 = model.fit(train_x, train_y)

m4 = mean_squared_error(test_y, model_4.predict(test_x))

df_4 = pd.Series(model_4.coef_, index=train_x.columns)
df_4['mean_squared_error'] = m4

fitted_models = [model_1, model_2, model_3, model_4]

return [df_1, df_2, df_3, df_4], fitted_models

```

```

In [45]: dfs, fitted_models= models(train_x, train_y)
outcome['RF'] = outcome['RF'].append(dfs[0].to_frame().T)
outcome['OLS'] = outcome['OLS'].append(dfs[1].to_frame().T)
outcome['RIDGE'] = outcome['RIDGE'].append(dfs[2].to_frame().T)
outcome['LASSO'] = outcome['LASSO'].append(dfs[3].to_frame().T)

```

## 5 5. Use bootstrap to get the variance of the estimators of the parameters

```

In [47]: # bootstrap in model
N = 100    # for getting the variance, 100-200 times paths are enough, as for percentiles
for i in range(N):
    x, y = resample(train_x, train_y)
    dfs, _ = models(x, y)
    outcome['RF'] = outcome['RF'].append(dfs[0].to_frame().T)
    outcome['OLS'] = outcome['OLS'].append(dfs[1].to_frame().T)
    outcome['RIDGE'] = outcome['RIDGE'].append(dfs[2].to_frame().T)
    outcome['LASSO'] = outcome['LASSO'].append(dfs[3].to_frame().T)

```

```

In [48]: results = {}
models_name = ['RF', 'OLS', 'RIDGE', 'LASSO']
for m in models_name:
    results[m] = pd.concat((outcome[m].iloc[0], outcome[m].iloc[1:].std()), 1)
    results[m].columns = [['features', 'std']]
    results[m] *= 1000
    results[m] = results[m].round(2)
result = pd.concat(results, 1)
result[['RF', 'OLS', 'RIDGE', 'LASSO']].to_csv('result.csv')

```

```
In [86]: result
```

```


Out[86]:
```

	LASSO	OLS	RF	RIDGE	\		
	features	std features	std features	std features			
OverallQual	108.84	9.12	113.90	9.37	168.99	11.95	107.08
YearBuilt	1.92	0.39	1.93	0.40	102.01	13.45	1.94

YearRemodAdd	3.45	0.45	3.36	0.46	47.97	8.98	3.43
ExterQual	0.00	4.59	16.37	14.85	52.21	13.07	10.74
BsmtQual	-10.16	7.19	-15.41	7.56	34.52	7.74	-15.30
TotalBsmtSF	0.24	0.03	0.24	0.03	77.01	10.82	0.24
1stFlrSF	0.06	0.04	0.06	0.04	61.35	6.83	0.06
GrLivArea	0.37	0.02	0.39	0.03	171.67	12.89	0.39
FullBath	-0.00	13.26	-34.32	18.83	44.48	12.19	-25.81
KitchenQual	-12.87	9.37	-25.71	9.39	25.47	8.80	-23.68
TotRmsAbvGrd	-3.07	6.09	-6.92	7.82	25.49	2.73	-7.31
GarageType	-10.46	4.18	-11.50	4.20	34.53	10.78	-11.87
GarageCars	0.00	7.95	26.57	21.51	60.95	11.17	20.01
GarageArea	0.32	0.04	0.23	0.06	65.66	10.38	0.26
MeanArea	0.12	0.03	0.12	0.03	27.69	3.90	0.12
mean_squared_error	107.45	3.92	104.55	4.15	51.54	1.79	105.90

	std
OverallQual	8.84
YearBuilt	0.39
YearRemodAdd	0.45
ExterQual	12.77
BsmtQual	7.18
TotalBsmtSF	0.03
1stFlrSF	0.04
GrLivArea	0.03
FullBath	14.95
KitchenQual	8.48
TotRmsAbvGrd	7.46
GarageType	4.15
GarageCars	15.73
GarageArea	0.05
MeanArea	0.03
mean_squared_error	4.02

```
In [85]: fig, axes = plt.subplots(2, 2, figsize=(10, 10))
for i, val in enumerate(models_name):
    axes[i//2, i%2].scatter(range(260), test_y, label="test_data")
    axes[i//2, i%2].scatter(range(260), fitted_models[i].predict(test_x), label="model")
    axes[i//2, i%2].set_ylabel('LogSalePrice', fontsize=13)
    axes[i//2, i%2].set_title(val)
    axes[i//2, i%2].legend()
plt.savefig("Comparison_of_test_data_with_predicted_data.jpg")
plt.show()
```



In [ ]: